

Container Native FS Interposer

Julia Hua, Jiawei Xiang, Hilario Gonzalez, Juncheng Cao

Mentors: Vasily Tarasov, Alex Merenstein

October 23, 2024

Recap

Container
Native FS
Interposer

Recap

Jaeger

FUSE

Plans

Last sprint

- Support CSI volume stacking
- CSI benchmark in CI
- Tracing FUSE filesystem

This sprint

- Deploy Jaeger into test cluster
- Flesh out Throttling FS
- Integrate Tracing FS with Jaeger
- Integrate Faulty fs with Jaeger

What is Jaeger

Container
Native FS
Interposer

Recap

Jaeger

FUSE

Plans

Jaeger

open source, distributed tracing platform

- monitor distributed workflows
- find & fix performance bottlenecks
- track down root causes
- analyze service dependencies

Span

Metric monitoring systems works on **Numbers**, but Jaeger operates on **Spans**

What is Span

Container
Native FS
Interposer

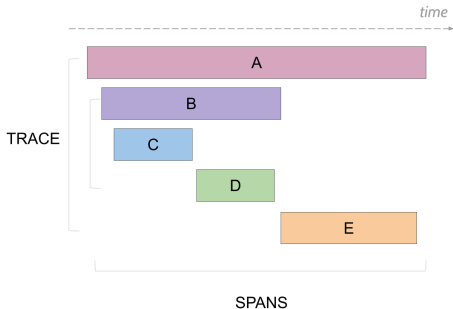

Recap

Jaeger

FUSE

Plans

*A **span** represents a **logical unit of work** that has an operation name, the start time of the operation, and the duration. Spans may be **nested** and ordered to model causal relationships.*



Jaeger Deployment Architecture


Container
Native FS
Interposer

Recap

Jaeger

FUSE

Plans



Workload Tracing

Container
Native FS
Interposer

Recap


Jaeger

FUSE

Plans

Progress

- Low-level API
- Export tracing data to Jaeger



Metric Collection

Container
Native FS
Interposer

Recap

Jaeger

FUSE

Plans

What

Gather quantitative and aggregate data on the performance and behavior of an application over time

- Frequency of each file operation
- Error rate
- Latency distribution

Why

- Understand type of workload
- Set up alerts when error rate or latency exceeds a threshold
- Monitor performance trends

IO Throttling

Container
Native FS
Interposer

Recap

Jaeger

FUSE

Plans

Problem

FUSE mandates IOs to be fully fulfilled (return exactly the number of bytes as requested)

Solution

Block the operation until we have enough IO quota left

Transition Faulty-fs Codebase to C++

Container
Native FS
Interposer

Recap

Jaeger

FUSE

Plans

Process

There were some simple changes that needed to be implemented widely throughout the file system code

- Dockerfile Compilation
- Included headers
- NULL → nullptr
- Change IO interface

OTEL for Simulated Error Logging

Container
Native FS
Interposer

Recap

Jaeger

FUSE

Plans

Process

- Configure OTEL SDK for application
- Instrument the error logging with OTEL's API
- Set up exporter to backend (Jaeger)

Why

- Integrate with OTEL used for tracing and metric collection
- Shared Jaeger backend

Future Plans

Container
Native FS
Interposer

Recap

Jaeger

FUSE

Plans

- Implement metric collection layer in FUSE
- Grafana dashboard for the metrics
- Make filesystem attributes configurable
- Combine all fuse filesystems into a single binary
- Post processing script for resolving file path
- Create parent spans covering multiple operations

Burndown

Container
Native FS
Interposer

Recap
Jaeger
FUSE
Plans

| Title | Date | Moved to Done Estimated Hours | Moved to Done Actual Hours | Ideal Hours | Burndown Estimated | Burndown Actual | Team # | Hours/Sprint |
|-------------------------------|------------|----------------------------------|-------------------------------|-------------|--------------------|-----------------|--------|--------------|
| | 2024/10/08 | 0 | 0 | 100 | 100 | 100 | 5 | 20 |
| | 2024/10/09 | 0 | 0 | 93 | 93 | 93 | | |
| Throttle IO with token bucket | 2024/10/10 | 10 | 9 | 86 | 76 | 77 | | |
| Deploy jaeger in test cluster | 2024/10/11 | 3 | 4 | 79 | 76 | 75 | | |
| Implement more faults | 2024/10/12 | 6 | 6 | 71 | 65 | 65 | | |
| Add test for token bucket | 2024/10/13 | 4 | 4 | 64 | 60 | 60 | | |
| Make tracing work with k8s | 2024/10/14 | 5 | 4 | 57 | 52 | 53 | | |
| Integrate throttling | 2024/10/15 | 3 | 3 | 50 | 47 | 47 | | |
| | 2024/10/16 | 0 | 0 | 43 | 43 | 43 | | |
| | 2024/10/17 | 0 | 0 | 36 | 36 | 36 | | |
| Fully implement tracing | 2024/10/18 | 5 | 5 | 29 | 24 | 24 | | |
| Implement more delay | 2024/10/19 | 4 | 5 | 21 | 17 | 16 | | |
| Modularize tracing | 2024/10/20 | 3 | 4 | 14 | 11 | 10 | | |
| | 2024/10/21 | 0 | 0 | 7 | 7 | 7 | | |
| Prepare for pre | 2024/10/22 | 6 | 6 | 0 | -6 | -6 | | |

Moved to Done Estimated Hours, Actual Hours, Ideal Hours, Running Total Estimate and Running Total Actual

