

Container Native FS Interposer

Julia Hua, Jiawei Xiang, Hilario Gonzalez, Juncheng Cao

Mentors: Vasily Tarasov, Alex Merenstein

November 20, 2024

Recap

Container
Native FS
Interposer

Recap

Workload

Configuration

Tracing

Faulty IO

Review

Last sprint

- Package and manage deployment with Helm/Helmfile
- Implement metric collection
- Integrate Faulty fs with Jaeger
- Fix TokenBucket implementation

This sprint

- Add postgresql and fio as test programs
- Implement metric collection FUSE layer
- Make FUSE layers individually configurable
- Implement tracing across multiple related operations
- Adding robustness to FaultyFs

PostgreSQL

Container
Native FS
Interposer

Recap

Workload

Configuration

Tracing

Faulty IO

Review

PostgreSQL

The World's Most Advanced Open Source Relational Database

- One of the most widely used database
- Good representative of real world workloads
- Builtin benchmarking tool: `pgbench`

FIO

Container
Native FS
Interposer

Recap

Workload

Configuration

Tracing

Faulty IO

Review

FIO

Flexible I/O tester

- Generic testing tool for generating various IO patterns
- Capable of recording and replaying IO traces
- A large collection of existing testing scripts

Limitations

IO tracing is done by blktrace at block level, not necessarily suitable for virtual filesystems

Configuring utilities

Container
Native FS
Interposer

Recap

Workload

Configuration

Tracing

Faulty IO

Review

```
{
  "metrics": {
    "enabled": true
  },
  "traces": {
    "enabled": true,
    "nestFileSpans": true,
    "otelLibName": "fstracing",
    "otelServiceName": "fs-workload-tracing",
    "otelHostName": "local-host",
    "otelEndpt": "localhost:4317"
  },
  "faultyIO": {
    "enabled": false
  },
  "throttleIO": {
    "enabled": false
  },
  "fakeIO": {
    "enabled": false
  }
}
```

Workload Tracing

Container
Native FS
Interposer

Recap

Workload

Configuration


Tracing

Faulty IO

Review

Feedback from mentors

■ Nesting spans



Adding robustness to FaultyFs

Container
Native FS
Interposer

Recap

Workload

Configuration

Tracing

Faulty IO

Review

More Testing

- Previously simulated faults in read, write, fsync, and get and set attributes
- Adding faults in certain functions can be unpredictable or breaks FS
- Some bugs with truncation indexing out of buffer
- Alternatively, more functions are well suited for faults that didn't have them before

More features

- Error log also gets file descriptor so more info available
- Implemented flags that can set the likelihood of faults and optionally increase when seeding happens

Review

Container
Native FS
Interposer

Recap

Workload

Configuration

Tracing

Faulty IO

Review

1	✓ Add prometheus-cpp and nlohmann_json dependency to container image #107	Nov 5, 2024
2	✓ Add postgresql/pgbench as a test program #108	Nov 6, 2024
3	✓ Add fio as a test program #109	Nov 6, 2024
4	✓ Generate FUSE config from volume attributes #110	Nov 7, 2024
5	✓ Address dependabot vulnerability report #113	Nov 7, 2024
6	✓ Post processing script for resolving file path from inode number #81	Nov 8, 2024
7	✓ Workload Tracing: Nest spans #120	Nov 10, 2024
8	✓ Make filesystem attributes configurable #83	Nov 11, 2024
9	✓ Create parent spans for multiple operations on the same file #80	Nov 11, 2024
10	✓ Implement metric collection layer for FUSE filesystem #117	Nov 12, 2024
11	✓ Make the FUSE layers individually configurable #118	Nov 14, 2024
12	✓ Prepare for presentation #119	Nov 18, 2024
13	🔄 Combine all fuse filesystems into a single binary #82	
14	🔄 Grafana dashboard for the metrics #84	
15	🔄 Test the installation in alternative kubernetes environments (like EKS) #87	
16	🔄 Push mapping from inode number to file name to external database #116	
17	🔄 OTLP in docker container for Faulty #121	

Burndown

Container
Native FS
Interposer

Recap

Workload

Configuration

Tracing

Faulty IO

Review

